

Infrastructure as a Service

Infrastructure as a Service, or IaaS is a service offering by most cloud providers that provides *virtual machines* and the associated infrastructure.

Virtual machines

In a hurry?

Modern CPUs have several operation modes:

- Ring 3 (unprivileged) runs the application
- Ring 0 runs the operating system kernel
- Ring -1 runs the hypervisor managing several kernels
- Ring -2 runs the Intel Management Engine

Other components are responsible for virtualizing other hardware components. For example, the IOMMU is responsible for virtualizing I/O.

Virtualization is a surprisingly old technology. The first virtualized system was the IBM System/370 mainframe with the VM/370.

When you think about mainframes you have to consider that these machines were *very* expensive and machine time was a precious commodity.

Initially CPUs in personal computers did not have application separation. The x86 line of Intel CPUs only received the *protected mode* in the 80386.

Protected mode introduced the concept of *rings* in the CPU. The operating system *kernel* would run in ring 0, device drivers in ring 1, and applications in ring 3.

Note

Device drivers today typically run on ring 0 instead of 1 or 2.

This ring system allowed the operating system to restrict the higher ring numbers from accessing certain functions or memory.

Note

If you try and set up a really old computer game like [Commander Keen](#) in [DOSBox](#) you will realize that you have to provide the game with a *virtual* memory space.

To work around the problems with protected mode the 80386 successor introduced [virtual mode](#). The new virtual 8086 mode allowed applications to run in protected mode.

For instance the CPU would create a simulated *virtual* memory space the program could write to and translate the virtual addresses to physical addresses.

Note

VM86 does not capture every instruction the application runs in virtual mode, only the sensitive CPU instructions. This enables legacy applications to run on modern CPUs.

In the mid 2000's CPUs became so powerful that it made sense to not only virtualize applications but whole operating systems. Some vendors modified kernel to facilitate them running in ring 3. Others employed [a number of techniques](#) we won't go into here.

Hardware vendors, of course, followed suit. In 2005 Intel added the VT-x (Vanderpool) feature to its new Pentium 4 CPUs for

VT-x and AMD-V added new ring -1 to accommodate *hypervisors*. This new ring allowed for separation between several op

Note

Intel also introduced a ring -2 for the [Intel Management Engine](#), a chip that functions as an OOBM in modern Intel chips. The ME

Virtualization also gave rise to Infrastructure as a Service. [AWS](#) was the first service that offered virtual machines as a service.

This allowed customers to create virtual machines as they needed it and they were billed for it on an hourly basis. (Later on

The presence of an API makes the difference between IaaS and plain old virtual machines as a service. IaaS allows a custo

What component of the software stack runs on Ring 3 in virtual mode?

- ☐ The application
- ☐ The kernel
- ☐ The hypervisor
- ☐ The management engine

What component of the software stack runs on Ring 0 in virtual mode?

- ☐ The application
- ☐ The kernel
- ☐ The hypervisor
- ☐ The management engine

What component of the software stack runs on Ring -1 in virtual mode?

- ☐ The application

- ☐ The kernel
- ☐ The hypervisor
- ☐ The management engine

What component of the software stack runs on Ring -2 in virtual mode?

- ☐ The application
- ☐ The kernel
- ☐ The hypervisor
- ☐ The management engine

What does virtualization mean?

- ☐ Every instruction by a virtual machine is captured by the kernel and translated.
- ☐ Critical instructions like memory operations are captured or translated by the kernel.
- ☐ Critical instructions like memory operations are captured or translated by the CPU and the hypervisor.

Typical instance types

When the cloud became popular in the late 2000s several providers attempted to offer a service that was fully dynamic in the sense that you could scale your resources up or down at will.

Instead most cloud providers nowadays opt to offer fixed machine sizes. To accommodate high-CPU and high RAM workloads, they offer different instance types.

- **Shared CPU:** These are small instances where a single CPU core is shared between multiple virtual machines, sometimes as many as 100.
- **Standard, dedicated core CPU:** These instance types receive one or more physical cores leading to a more stable performance.
- **High CPU:** These instance types are usually hosted on physical servers that have a very high CPU to RAM ratio. According to AWS, the ratio is 1:1 for these instances.
- **High RAM:** This offering is the exact opposite of the high CPU offering. The machines on offer here include more RAM than CPU.
- **Storage:** These instance types contain large amounts of local storage (see below in the storage section).
- **Hardware-specific:** These instance types offer access to dedicated hardware such as graphics cards (GPUs) or FPGAs.

Automation

In a hurry?

- Cloud-init allows for running a script, or other initial configuration on virtual machines on first boot.
- It is also responsible for managing password resets when desired. It can be used to fully automate the setup of a virtual machine.
- Terraform and Ansible are tools that interact with the cloud API to provision virtual machines programmatically.
- Ansible is also capable of running inside a virtual machine to configure the software within.
- Terraform requires full control of the machines it is managing and implements what's called immutable infrastructure.

As discussed before, what makes an IaaS cloud provider a cloud provider is the fact that they offer an API to automate the process.

Initially this problem would be solved by creating *templates* for the operating system that launches. In larger cloud setups this is not practical.

Thankfully in the last decade a lot has happened and [Cloud Init](#) has established itself as a defacto standard in the IaaS world.

A DevOps engineer can simply inject a script that runs at the first start that takes care of all the installation steps required.

Tools like [Terraform](#) or [Ansible](#) assist with managing the whole process of provisioning the virtual machines and supplying it with user data.

What is the role of cloud-init?

- ☐ It initializes a cloud account.
- ☐ It creates a virtual machine.
- ☐ It runs initial machine configuration on a virtual machine.

Virtual machine pools

In a hurry?

- Virtual machine pools automatically create and destroy machines to keep up a desired pool size.
- Some implementations also have autoscaling.

One other use of user data are virtual machine pools. Each cloud provider adopts a different name for them, ranging from *auto-scaling groups* to *elastic pools*.

The number of machines in a pool can, of course, be changed either manually or in some cases automatically using rules for scaling.

Combined with the aforementioned user data this can be a very powerful tool to create a dynamically sized pool of machines.

These pools are often integrated with the various load-balancer offerings cloud providers have in their portfolio to direct traffic to the machines.

Storage

In a hurry?

- Local disks offer affordable performance at the cost of redundancy.
- Network block storage offers resilience to machine failures, but costs more to ensure the same performance. Not all NBS implements RAID.
- Network file systems offer access from multiple virtual machines in parallel at the cost of performance.
- Object storage offers parallel access from multiple VMs and scalability at the cost of performance and consistency.
- Object storages are typically integrated on the application level rather than the OS level.

When it comes to data storage virtual machines work exactly like your physical machine would: there is a physical disk (or more).

However, a distributed storage system is generally either slower or more expensive for the same performance by several magnitudes.

When we talk about storage systems we are talking about two types: block devices and filesystems. On the physical disk data is stored in blocks, that's something the filesystem must keep track of.

Therefore we traditionally call raw disk devices *block devices*. Block devices are (with very few exceptions) only accessible by a single VM.

Local Storage

As described above the simplest and most widely supported option to store data from your virtual machine is a disk that is located on the physical machine.

Some cloud providers offer disk redundancy ([RAID](#)) while others don't. At any rate a hardware failure of the physical machine is possible.

It is therefore very advisable to solve redundancy on top of the virtual machine, e.g. by building a replicated database setup.

Which of the following is provided by local storage?

- ☐ Fault-tolerance in the face of a machine failure.
- ☐ High IO performance.
- ☐ The ability to move the data volume to a different machine.
- ☐ The ability to access the data volume from several machines at once.
- ☐ Data consistency.

Network Block Storage

Network block storage means a block storage that is delivered over the network. The network here can mean a traditional IP network or a specialized storage network.

As described before block storage is, in general, single-VM only. You can't access the files stored on a block storage device from multiple VMs.

Also note that Network Block Storage does not automatically come with redundancy. Some solutions, such as [iSCSI](#) simply

At any rate, using Network Block Storage does not absolve you from the duty to make backups and have a documented an

Which of the following is provided by network block storage?

- ☐ Fault-tolerance in the face of a machine failure.
- ☐ High IO performance.
- ☐ The ability to move the data volume to a different machine.
- ☐ The ability to access the data volume from several machines at once.
- ☐ Data consistency.

Network File Systems

In contrast to network block storage network file systems offer access to data not on a block level, but on a file level. Over th

The filesystem has to keep track of which machine has which file open, or has locks on which file. When machine edit the s

Which of the following is provided by network filesystems?

- ☐ Fault-tolerance in the face of a machine failure.
- ☐ High IO performance.
- ☐ The ability to move the data volume to a different machine.
- ☐ The ability to access the data volume from several machines at once.
- ☐ Data consistency.

Object storage

Object storage systems are similar to network file systems in that they deal with files rather than blocks. However, they do n

While object storages technically *can* be used as a filesystem on an operating system level for example by using [s3fs](#) this is

Operating system level integration should only be used as a last resort and object storages should be integrated on the app

Which of the following is provided by object storages?

- ☐ Fault-tolerance in the face of a machine failure.
- ☐ High IO performance.
- ☐ The ability to move the data volume to a different machine.
- ☐ The ability to access the data volume from several machines at once.
- ☐ Data consistency.

Which storage type is Amazon's EBS?

- ☐ Local disk
- ☐ Network block storage
- ☐ Network filesystem
- ☐ Object storage

Which storage type is Amazon's EFS?

- ☐ Local disk
- ☐ Network block storage
- ☐ Network filesystem
- ☐ Object storage

Which storage type is Ceph RBD?

- ☐ Local disk

- ☐ Network block storage
- ☐ Network filesystem
- ☐ Object storage

Which storage type is iSCSI?

- ☐ Local disk
- ☐ Network block storage
- ☐ Network filesystem
- ☐ Object storage

Which storage type is S3?

- ☐ Local disk
- ☐ Network block storage
- ☐ Network filesystem
- ☐ Object storage

Network

The next big topic concerning IaaS services is networks. Before we go into the cloud-aspect let's look at how the underlying

How cloud networks are built

So, let's get started. Imagine a data center from the first lecture. Your task is to build an IaaS cloud provider. You put your servers

This sounds like a lot of bandwidth available but keep in mind that your virtual machines get assigned to the physical machines


This is part of the reason why in the cloud scaling horizontally (adding more machines) is preferred rather than creating huge

Network architectures offered by cloud providers

When we look at the network offerings by cloud providers there are three types:

1. **Private-only network with NAT:** This option is provided by the larger cloud providers such as [AWS](#), [Azure](#), [GCP](#) and IaaS virtual machine and they can work independently.

2. **Default public IP:** This option is provided by smaller IaaS providers such as [DigitalOcean](#), [Exoscale](#), [Hetzner](#), [Linode](#),



3. **Fully customizable:** This setup allows the customer to design their network connectivity as they see fit. This setup is s



Start Center

Test x

Palette



Server



HDD
Storage



SSD
Storage



CD ROM



Internet



Load
Balancer



Composite
Instances



Note

There are several other cloud providers which we have no information on, such as the [Deutsche Telekom/Open Telekom Cloud](#),

Note

Out of group 2 it is worth mentioning that the services that are available on the public network (firewalls, load balancers) are often

Firewalling

IaaS providers often also offer network firewalls as a service, included in the platform. Firewalls generally have two rule types

Firewall providers often employ the concept of *security groups*. The implementation varies greatly, but in general security groups

For most cloud providers you will need to create an explicit rule allowing traffic to flow between two machines in the same security group

The advantage of security groups is that the rules can be made in such a way that they reference other security groups rather than

Add Rule to backend

Type *

INGRESS

Protocol *

TCP

UDP

ICMP

IC

Source Type *

CIDR

Source *

frontend

Port(s) *

3306

Insert a range of ports you wish to expose.
As an example, for an incoming HTTP connection

Description

ADD

What do security groups offer?

- ☐ Filtering based on IP address
- ☐ Filtering based on the requested service
- ☐ Filtering based on the requested domain name

- ☐ Filtering based on the requested subpage on a website

Network load balancers

Network load balancers are an option some cloud providers offer. In contrast to Application Load Balancers they do not offer

Depending on the cloud provider in question network load balancers may or may not offer terminating encrypted connections.

When designing an architecture it is worth considering if the real IP address of the connecting client will be needed. If the backend needs the real IP address, a load balancer that terminates the connection and forwards the request to the backend will not work.

In order to make sure requests are not sent to faulty backends NLBs include a health check feature. This health check either sends requests to the backends or probes the backends. In either case, the health check can detect faulty machines.

When talking about load balancers an interesting question is the load balancing strategy. Most load balancers support either round-robin or weighted round-robin.

What do NLBs typically offer?

- ☐ Spreading incoming connections across multiple backend machines equally.
- ☐ Spreading incoming connections across multiple backend machines, sending connections from the same client to the same backend machine.
- ☐ Spreading incoming connections across multiple backend machines, based on the domain name of the request.
- ☐ Spreading incoming connections across multiple backend machines, based on the subpage requested.
- ☐ Terminating encrypted connections so the backend doesn't have to (SSL/TLS offloading).

VPNs, private interconnects, and routing services

While it seems convenient at first to use only the public network several organizations have security models that prevent access to the public network.

However, this presents a problem when moving data between several, geographically distributed locations. Most companies use MPLS tunnels to solve this problem.

MPLS tunnels create a virtual connectivity that does not go on the Internet. While being expensive and slow to set up, it can be used to connect several locations.

VPN's on the other hand create a virtual connectivity by sending data over the Internet in an encrypted form. Bandwidth or latency can be a problem.

Larger cloud providers tend to offer both options. MPLS is supported by the larger cloud providers ([AWS Direct connect](#), [Azure ExpressRoute](#), or [Google Cloud Interconnect](#)).

VPN is also offered mostly by large providers ([AWS VPC VPN](#), [Azure VPN](#), or [Google Cloud VPN](#)). However, keep in mind that not all providers support the same features.

It is also worth noting that VPN's can be used to connect cloud providers together.

What VPN type is offered by all major cloud providers?

- ☐ Site-to-site
- ☐ Device-to-site
- ☐ Device-to-device

What VPN protocol is offered by all major cloud providers?

- ☐ OpenVPN
- ☐ IPsec
- ☐ SSTP
- ☐ L2TP
- ☐ PPTP

What VPN type can IPsec offer by itself?

- ☐ Site-to-site
- ☐ Device-to-site
- ☐ Device-to-device

DNS

The [Domain Name Service](#) is one of the services that are all but required for building an infrastructure. It provides domain n

There is a difference, however, between DNS services on offer. Some DNS services by cloud providers offer only simple res

More advanced features may include automatic DNS failover. This involves running regular health checks on your services

Monitoring

Some cloud providers offer includes basic monitoring, such as CPU or memory usage. Some providers are offering monitor

Often times the monitoring facilities offered by cloud providers are not sufficient for keeping an application running and more