

Infrastructure as a Service

Infrastructure as a Service, or IaaS is a service offering by most cloud providers that provides *virtual machines* and the associated infrastructure.

Virtual machines

Virtualization is a surprisingly old technology. The first virtualized system was the IBM System/370 mainframe with the VM/370.

When you think about mainframes you have to consider that these machines were *very* expensive and machine time was a precious commodity.

Initially CPUs in personal computers did not have application separation. The x86 line of Intel CPUs only received the *protected mode*.

Protected mode introduced the concept of *rings* in the CPU. The operating system *kernel* would run in ring 0, device drivers in ring 1, and applications in ring 3.

Note

Device drivers today typically run on ring 0 instead of 1 or 2.

This ring system allowed the operating system to restrict the higher ring numbers from accessing certain functions or memory.

Note

If you try and set up a really old computer game like [Commander Keen](#) in [DOSBox](#) you will realize that you have to provide the game with a virtualized hardware environment.

To work around the problems with protected mode the 80386 successor introduced [virtual mode](#). The new virtual 8086 mode allowed applications to run in protected mode.

For instance the CPU would create a simulated *virtual* memory space the program could write to and translate the virtual addresses to physical addresses.

Note

VM86 does not capture every instruction the application runs in virtual mode, only the sensitive CPU instructions. This enables legacy applications to run on modern CPUs.

In the mid 2000's CPUs became so powerful that it made sense to not only virtualize applications but whole operating systems. Some vendors modified the kernel to facilitate them running in ring 3. Others employed [a number of techniques](#) we won't go into here.

Hardware vendors, of course, followed suit. In 2005 Intel added the VT-x (Vanderpool) feature to its new Pentium 4 CPUs for virtualization.

VT-x and AMD-V added new ring -1 to accommodate *hypervisors*. This new ring allowed for separation between several operating systems running on the same hardware.

Note

Intel also introduced a ring -2 for the [Intel Management Engine](#), a chip that functions as an OOBM in modern Intel chips. The ME is used for various management functions.

Virtualization also gave rise to Infrastructure as a Service. [AWS](#) was the first service that offered virtual machines as a service.

This allowed customers to create virtual machines as they needed it and they were billed for it on an hourly basis. (Later on

The presence of an API makes the difference between IaaS and plain old virtual machines as a service. IaaS allows a custo

Typical instance types

When the cloud became popular in the late 2000s several providers attempted to offer a service that was fully dynamic in th

Instead most cloud providers nowadays opt to offer fixed machine sizes. To accommodate high-CPU and high RAM workloa

- **Shared CPU:** These are small instances where a single CPU core is shared between multiple virtual machines, somet
- **Standard, dedicated core CPU:** These instance types receive one or more physical cores leading to a more stable pe
- **High CPU:** These instance types are usually hosted on physical servers that have a very high CPU to RAM ratio. Accor
- **High RAM:** This offering is the exact opposite of the high CPU offering. The machines on offer here include more RAM
- **Storage:** These instance types contain large amounts of local storage (see below in the storage section).
- **Hardware-specific:** These instance types offer access to dedicated hardware such as graphics cards (GPUs) or FPGA

Automation

As discussed before, that makes an IaaS cloud provider a cloud provider is the fact that they offer an API to automate the p

Initially this problem would be solved by creating *templates* for the operating system that launches. In larger cloud setups th

Thankfully in the last decade a lot has happened and [Cloud Init](#) has established itself as a defacto standard in the IaaS worl

A DevOps engineer can simply inject a script that runs at the first start that takes care of all the installation steps required.

Tools like [Terraform](#) or [Ansible](#) assist with managing the whole process of provisioning the virtual machines and supplying it

Virtual machine pools

One other use of user data are virtual machine pools. Each cloud provider adopts a different name for them, ranging from in

The number of machines in a pool can, of course, be changed either manually or in some cases automatically using rules fo

Combined with the aforementioned user data this can be a very powerful tool to create a dynamically sized pool of machine

These pools are often integrated with the various load-balancer offerings cloud providers have in their portfolio to direct traff

Storage

When it comes to data storage virtual machines work exactly like your physical machine would: there is a physical disk (or m

However, a distributed storage system is generally either slower or more expensive for the same performance by several m

When we talk about storage systems we are talking about two types: block devices and filesystems. On they physical disk c
so that's something the filesystem must keep track of.

Therefore we traditionally call raw disk devices *block devices*. Block devices are (with very few exceptions) only accessible
can only ever be used by a single VM.

Local Storage

As described above the simplest and most widely supported option to store data from your virtual machine is a disk that is local to the physical machine. Some cloud providers offer disk redundancy ([RAID](#)) while others don't. At any rate a hardware failure on the physical machine is a problem. It is therefore very advisable to solve redundancy on top of the virtual machine, e.g. by building a replicated database setup.

Network Block Storage

Network block storage means a block storage that is delivered over the network. The network here can mean a traditional IP network. As described before block storage is, in general, single-VM only. You can't access the files stored on a block storage device from multiple VMs. Also note that Network Block Storage does not automatically come with redundancy. Some solutions, such as [iSCSI](#) simply replicate the data. At any rate, using Network Block Storage does not absolve you from the duty to make backups and have a documented and tested recovery plan.

Network File Systems

In contrast to network block storage network file systems offer access to data not on a block level, but on a file level. Over the network, the filesystem has to keep track of which machine has which file open, or has locks on which file. When a machine edits the file, the filesystem has to ensure that the data is consistent.

Object storage

Object storage systems are similar to network file systems in that they deal with files rather than blocks. However, they do not have a filesystem. While object storages technically *can* be used as a filesystem on an operating system level for example by using [s3fs](#) this is not recommended. Operating system level integration should only be used as a last resort and object storages should be integrated on the application level.

Network

The next big topic concerning IaaS services is networks. Before we go into the cloud-aspect let's look at how the underlying network is built.

How cloud networks are built

So, let's get started. Imagine a data center from the first lecture. Your task is to build an IaaS cloud provider. You put your servers in a data center. This sounds like a lot of bandwidth available but keep in mind that your virtual machines get assigned to the physical machines. This is part of the reason why in the cloud scaling horizontally (adding more machines) is preferred rather than creating huge virtual machines.

Network architectures offered by cloud providers

When we look at the network offerings by cloud providers there are three types:

1. **Private-only network with NAT:** This option is provided by the larger cloud providers such as [AWS](#), [Azure](#), [GCP](#) and [IBM Cloud](#). Each virtual machine has a private IP address and they can work independently.
2. **Default public IP:** This option is provided by smaller IaaS providers such as [DigitalOcean](#), [Exoscale](#), [Hetzner](#), [Linode](#), [Vultr](#) and [Gigamonks](#).

3. **Fully customizable:** This setup allows the customer to design their network connectivity as they see fit. This setup is s



Start Center

Test ×

▸ Palette



Server



HDD
Storage



SSD
Storage



CD ROM




Internet



Load
Balancer



Composite
Instances



Note

There are several other cloud providers which we have no information on, such as the [Deutsche/Open Telekom Cloud](#), or the [Ali](#)

Note

Out of group 2 it is worth mentioning that the services that are available on the public network (firewalls, load balancers) are ofte

TODO: add illustration

Firewalling

IaaS providers often also offer network firewalls as a service, included in the platform. Firewalls generally have two rule types

Firewall providers often employ the concept of *security groups*. The implementation varies greatly, but in general security groups

For most cloud providers you will need to create an explicit rule allowing traffic to flow between two machines in the same s

The advantage of security groups is that the rules can be made in such a way that they reference other security groups rath

TODO: add screenshot of security group configuration

Network load balancers

Network load balancers are an option some cloud providers offer. In contrast to Application Load Balancers they do not offe

TODO add illustration

Depending on the cloud provider in question network load balancers may or may not offer terminating encrypted connection

When designing an architecture it is worth considering if the real IP address of the connecting client will be needed. If the ba

When talking about load balancers an interesting question is the load balancing strategy. Most load balancers support eithe

VPNs, private interconnects, and routing services

While it seems convenient at first to use only the public network several organizations have security models that prevent ac

However, this presents a problem when moving data between several, geographically distributed locations. Most companies

MPLS tunnels create a virtual connectivity that does not go on the Internet. While being expensive and slow to set up, it can

VPN's on the other hand create a virtual connectivity by sending data over the Internet in an encrypted form. Bandwidth or l

Larger cloud providers tend to offer both options. MPLS is supported by the larger cloud providers ([AWS Direct connect](#), [Az](#)

VPN is also offered mostly by large providers ([AWS VPC VPN](#), [Azure VPN](#), or [Google Cloud VPN](#)). However, keep in mind

It is also worth noting that VPN's can be used to connect cloud providers together.

DNS

The [Domain Name Service](#) is one of the services that are all but required for building an infrastructure. It provides domain n

There is a difference, however, between DNS services on offer. Some DNS services by cloud providers offer only simple res

More advanced features may include automatic DNS failover. This involves running regular health checks on your services

Monitoring

Some cloud providers offer included basic monitoring, such as CPU or memory usage. Some providers are offering monitor

Automation